

▪ [Reference](#)

```
import numpy as np
import PIL
import os
from glob import glob
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow.keras import layers
import cv2
import random
from keras.utils.vis_utils import plot_model
```

```
MAIN_PATH = "../input/celeba-dataset/img_align_celeba/img_align_celeba"
```

```
image_paths = glob(MAIN_PATH+"/*")
```


```
len(image_paths)
```

202599

```
def readImage(path, image_size=(256,256)):
    img = np.asarray(PIL.Image.open(path).resize(image_size))
    img = ((img - 127.5) / 127.5).astype("float32")
    return img
```

```
test_img = readImage("../input/celeba-dataset/img_align_celeba/img_align_celeba/000030.jpg")
print(test_img.shape)
```

```
plt.imshow(test_img)
plt.show()
```



- I said we'll read images as batches so now we'll define a generator that yield one batch data.

```
BATCH_SIZE = 128
STEPS_PER_EPOCH = 500
print("Steps per epochh are",STEPS_PER_EPOCH)
def dataGenerator(batch_size):
    while True:
        paths = random.choices(image_paths,k=batch_size)
        batch = []
        for p in paths:
            batch.append(readImage(p))

        yield np.asarray(batch)
```

```
dataGen = dataGenerator(BATCH_SIZE)
print(next(dataGen).shape)
```

Steps per epochh are 500
(128, 256, 256, 3)

```
WEIGHT_INIT = tf.keras.initializers.RandomNormal(mean=0.0,stddev=0.2)
def make_generator():
    model = tf.keras.Sequential()

    model.add(layers.Dense(16*16*256,use_bias=False,input_shape=(100,)))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())
    model.add(layers.Reshape((16,16,256)))

    assert model.output_shape == (None,16,16,256)

    model.add(layers.Conv2DTranspose(128,(5,5),strides=(2,2),use_bias=False,padding="same",kernel_initializer=WEIGHT_INIT))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    assert model.output_shape == (None,32,32,128)

    model.add(layers.Conv2DTranspose(128,(5,5),strides=(2,2),use_bias=False,padding="same",kernel_initializer=WEIGHT_INIT))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    assert model.output_shape == (None,64,64,128)

    model.add(layers.Conv2DTranspose(64,(5,5),strides=(2,2),use_bias=False,padding="same",kernel_initializer=WEIGHT_INIT))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    assert model.output_shape == (None,128,128,64)

    model.add(layers.Conv2DTranspose(3,(5,5),strides=(2,2),use_bias=False,padding="same",kernel_initializer=WEIGHT_INIT,
        activation="tanh"
    ))
    assert model.output_shape == (None,256,256,3)
    return model
```

```
generator = make_generator()
```

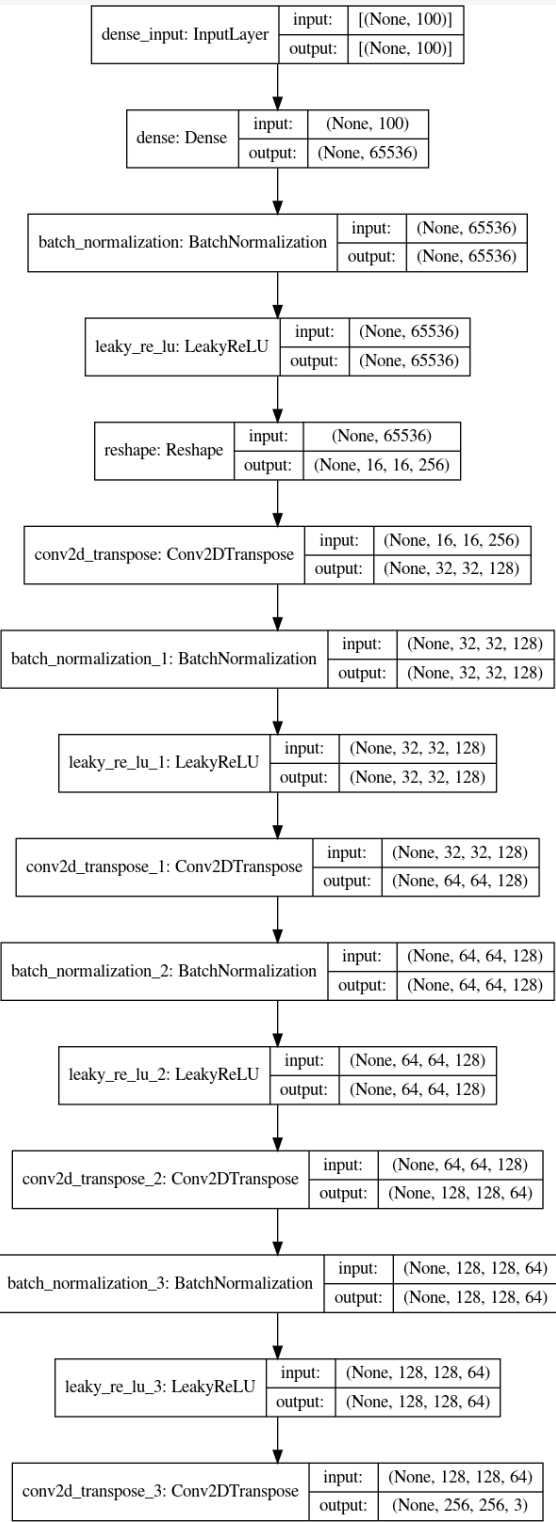
```
generator.summary()
```

Model: "sequential"		
Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 65536)	6553600

batch_normalization (BatchNo	(None, 65536)	262144

leaky_re_lu (LeakyReLU)	(None, 65536)	0
reshape (Reshape)	(None, 16, 16, 256)	0
conv2d_transpose (Conv2DTran	(None, 32, 32, 128)	819200
batch_normalization_1 (Batch	(None, 32, 32, 128)	512
leaky_re_lu_1 (LeakyReLU)	(None, 32, 32, 128)	0
conv2d_transpose_1 (Conv2DTr	(None, 64, 64, 128)	409600
batch_normalization_2 (Batch	(None, 64, 64, 128)	512
leaky_re_lu_2 (LeakyReLU)	(None, 64, 64, 128)	0
conv2d_transpose_2 (Conv2DTr	(None, 128, 128, 64)	204800
batch_normalization_3 (Batch	(None, 128, 128, 64)	256
leaky_re_lu_3 (LeakyReLU)	(None, 128, 128, 64)	0
conv2d_transpose_3 (Conv2DTr	(None, 256, 256, 3)	4800
=====		
Total params: 8,255,424		
Trainable params: 8,123,712		
Non-trainable params: 131,712		

```
plot_model(generator,show_shapes = True, show_layer_names = True, to_file='Generator_Model.png')
```



```
cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)

def generator_loss(fake_output):
    return cross_entropy(tf.ones_like(fake_output),fake_output)

gen_optimizer = tf.keras.optimizers.Adam(lr=1e-4)
```

```
def make_discriminator():
    model = tf.keras.Sequential()

    model.add(layers.Conv2D(64, (5,5), strides=(2,2), padding="same", input_shape=(256, 256, 3)))
    model.add(layers.LeakyReLU())
    model.add(layers.Dropout(0.3))

    model.add(layers.Conv2D(128, (5,5), strides=(2,2), padding="same"))
    model.add(layers.LeakyReLU())
    model.add(layers.Dropout(0.3))

    model.add(layers.Conv2D(256, (5,5), strides=(2,2), padding="same"))
    model.add(layers.LeakyReLU())
    model.add(layers.Dropout(0.3))

    model.add(layers.Flatten())
    model.add(layers.Dense(1))

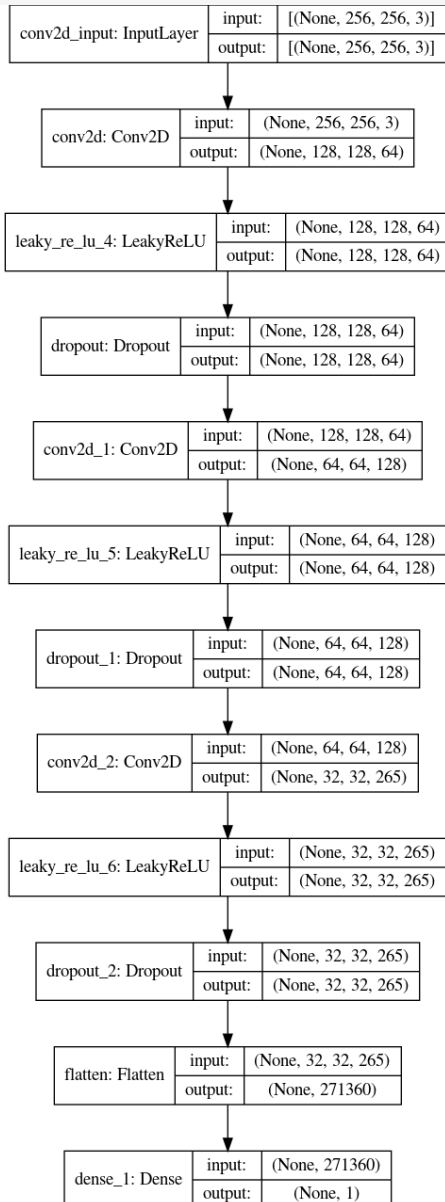
    return model
```

```
discriminator = make_discriminator()
```

```
discriminator.summary()
```

Model: "sequential_1"		
Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 128, 128, 64)	4864
leaky_re_lu_4 (LeakyReLU)	(None, 128, 128, 64)	0
dropout (Dropout)	(None, 128, 128, 64)	0
conv2d_1 (Conv2D)	(None, 64, 64, 128)	204928
leaky_re_lu_5 (LeakyReLU)	(None, 64, 64, 128)	0
dropout_1 (Dropout)	(None, 64, 64, 128)	0
conv2d_2 (Conv2D)	(None, 32, 32, 256)	848265
leaky_re_lu_6 (LeakyReLU)	(None, 32, 32, 256)	0
dropout_2 (Dropout)	(None, 32, 32, 256)	0
Flatten (Flatten)	(None, 271360)	0
dense_1 (Dense)	(None, 1)	271361
=====		
Total params: 1,329,418		
Trainable params: 1,329,418		
Non-trainable params: 0		

```
plot_model(discriminator, show_shapes = True, show_layer_names = True, to_file='Generator_Model.png')
```



```
def discriminator_loss(real_images, fake_images):
    real_loss = cross_entropy(tf.ones_like(real_images), real_images)
    fake_loss = cross_entropy(tf.zeros_like(fake_images), fake_images)
    total_loss = real_loss + fake_loss
    return total_loss
```

```
discriminator_optimizer = tf.keras.optimizers.Adam(lr=1e-4)
```

```
EPOCHS = 5
NOISE_DIM = 100
```

```
@tf.function
def train_step(images):
    noise = tf.random.normal([BATCH_SIZE, NOISE_DIM])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

        real_output = discriminator(images, training=True)
        fake_output = discriminator(generated_images, training=True)
        L_G = tf.reduce_mean(tf.abs(real_output - fake_output))
        L_D = tf.reduce_mean(tf.abs(real_output - fake_output)) - k * tf.reduce_mean(tf.abs(real_output - fake_output))

        gen_loss = generator_loss(fake_output)
        disc_loss = discriminator_loss(real_output, fake_output)

        gradients_of_generator = gen_tape.gradient(L_G, generator.trainable_variables)
        gradients_of_discriminator = disc_tape.gradient(L_D, discriminator.trainable_variables)

        gen_optimizer.apply_gradients(zip(gradients_of_generator, generator.trainable_variables))
        discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator, discriminator.trainable_variables))
```

```
import time
import sys
def train(epochs):
    for epoch in range(epochs):
        start = time.time()
        for step in range(STEPS_PER_EPOCH):
            train_step(next(dataGen))

            sys.stdout.write(f"\rSTEP: {step}/{STEPS_PER_EPOCH}")
            sys.stdout.flush()

        finish_time = round(time.time() - start, 2)
        print(f"Epoch {epoch}/{epochs} Process Time : {finish_time}")
        print("-"*15)
```

```
train(EPOCHS)

STEP: 499/500Epoch 0/5 Process Time : 640.75
-----
STEP: 499/500Epoch 1/5 Process Time : 501.48
-----
STEP: 499/500Epoch 2/5 Process Time : 485.45
-----
STEP: 499/500Epoch 3/5 Process Time : 444.74
-----
STEP: 499/500Epoch 4/5 Process Time : 423.32
-----
```

```
train(3)

STEP: 499/500Epoch 0/3 Process Time : 417.55
-----
STEP: 499/500Epoch 1/3 Process Time : 415.64
-----
STEP: 499/500Epoch 2/3 Process Time : 414.8
-----
```


```
train(2)

STEP: 499/500Epoch 0/2 Process Time : 414.82
-----
STEP: 499/500Epoch 1/2 Process Time : 414.91
-----
```

```
noise = tf.random.normal([16, 100])
generated_images = np.asarray(generator(noise, training=False))

fig = plt.figure(figsize=(6, 6))
for i in range(16):
    plt.subplot(4, 4, i+1)
    plt.imshow((generated_images[i, :, :] * 127.5 + 127.5).astype("int"))
    plt.axis("off")

plt.show()
```



Compared to the results presented in the BEGAN paper, the images generated by my model are clearly lacking in quality and coherence. While the BEGAN model was able to produce high-resolution (up to 128x128 pixels), diverse, and natural-looking face images, my model is struggling to generate anything beyond blurred and disoriented faces.